

PROBLEM STATEMENT:-

TO PREDICT THE RAIN FALL BASED ON VARIOUS FEATURES OF THE DATASET

1.Data Collection

```
In [1]: #importing libraries
import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn import preprocessing,svm
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv(r"C:\Users\DELL\Downloads\rainfall in india 1901-2015.csv")
df
```

Out[2]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan- Feb	Mar- May	Jun- Sep
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.2	33.6	3373.2	136.3	560.3	1696.3
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.0	160.5	3520.7	159.8	458.3	2185.9
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.4	225.0	2957.4	156.7	236.1	1874.0
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.7	40.1	3079.6	24.1	506.9	1977.6
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25.4	344.7	2566.7	1.3	309.7	1624.9
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	184.3	14.9	1533.7	7.9	196.2	1013.0
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	12.4	8.8	1405.5	19.3	99.6	1119.5
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	78.1	26.7	1426.3	60.6	131.1	1057.0
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2	59.0	62.3	1395.0	69.3	76.7	958.5
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	231.0	159.0	1642.9	2.7	223.9	860.9

4116 rows × 19 columns

## 2.Data Preprocessing

In [3]: df.head()

Out[3]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	Jun-Sep	Oct-Dec
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.2	33.6	3373.2	136.3	560.3	1696.3	980.3
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.0	160.5	3520.7	159.8	458.3	2185.9	716.7
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.4	225.0	2957.4	156.7	236.1	1874.0	690.6
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.7	40.1	3079.6	24.1	506.9	1977.6	571.0
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25.4	344.7	2566.7	1.3	309.7	1624.9	630.8

In [4]: df.tail()

Out[4]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	Jun-Sep	Oct-Dec
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	184.3	14.9	1533.7	7.9	196.2	1013.0	316
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	12.4	8.8	1405.5	19.3	99.6	1119.5	167
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	78.1	26.7	1426.3	60.6	131.1	1057.0	177
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2	59.0	62.3	1395.0	69.3	76.7	958.5	290
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	231.0	159.0	1642.9	2.7	223.9	860.9	555

In [5]: df.describe()

Out[5]:

	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP
count	4116.000000	4112.000000	4113.000000	4110.000000	4112.000000	4113.000000	4111.000000	4109.000000	4112.000000	4110.000000
mean	1958.218659	18.957320	21.805325	27.359197	43.127432	85.745417	230.234444	347.214334	290.263497	197.361922
std	33.140898	33.585371	35.909488	46.959424	67.831168	123.234904	234.710758	269.539667	188.770477	135.408345
min	1901.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.400000	0.000000	0.000000	0.100000
25%	1930.000000	0.600000	0.600000	1.000000	3.000000	8.600000	70.350000	175.600000	155.975000	100.525000
50%	1958.000000	6.000000	6.700000	7.800000	15.700000	36.600000	138.700000	284.800000	259.400000	173.900000
75%	1987.000000	22.200000	26.800000	31.300000	49.950000	97.200000	305.150000	418.400000	377.800000	265.800000
max	2015.000000	583.700000	403.500000	605.600000	595.100000	1168.600000	1609.900000	2362.800000	1664.600000	1222.000000

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4116 entries, 0 to 4115
Data columns (total 19 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   SUBDIVISION     4116 non-null   object  
1   YEAR            4116 non-null   int64   
2   JAN             4112 non-null   float64  
3   FEB             4113 non-null   float64  
4   MAR             4110 non-null   float64  
5   APR             4112 non-null   float64  
6   MAY             4113 non-null   float64  
7   JUN             4111 non-null   float64  
8   JUL             4109 non-null   float64  
9   AUG             4112 non-null   float64  
10  SEP             4110 non-null   float64  
11  OCT             4109 non-null   float64  
12  NOV             4105 non-null   float64  
13  DEC             4106 non-null   float64  
14  ANNUAL          4090 non-null   float64  
15  Jan-Feb        4110 non-null   float64  
16  Mar-May        4107 non-null   float64  
17  Jun-Sep        4106 non-null   float64  
18  Oct-Dec        4103 non-null   float64  
dtypes: float64(17), int64(1), object(1)
memory usage: 611.1+ KB
```

```
In [7]: df.isnull().sum()
```

```
Out[7]: SUBDIVISION     0
YEAR                 0
JAN                  4
FEB                  3
MAR                  6
APR                  4
MAY                  3
JUN                  5
JUL                  7
AUG                  4
SEP                  6
OCT                  7
NOV                 11
DEC                 10
ANNUAL              26
Jan-Feb             6
Mar-May             9
Jun-Sep            10
Oct-Dec            13
dtype: int64
```

```
In [10]: df.fillna(method='ffill',inplace=True)
```

```
In [11]: df.isnull().sum()
```

```
Out[11]: SUBDIVISION    0
YEAR                0
JAN                 0
FEB                 0
MAR                 0
APR                 0
MAY                 0
JUN                 0
JUL                 0
AUG                 0
SEP                 0
OCT                 0
NOV                 0
DEC                 0
ANNUAL              0
Jan-Feb             0
Mar-May             0
Jun-Sep             0
Oct-Dec             0
dtype: int64
```

```
In [12]: df.columns
```

```
Out[12]: Index(['SUBDIVISION', 'YEAR', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL',
               'AUG', 'SEP', 'OCT', 'NOV', 'DEC', 'ANNUAL', 'Jan-Feb', 'Mar-May',
               'Jun-Sep', 'Oct-Dec'],
              dtype='object')
```

```
In [13]: df.shape
```

```
Out[13]: (4116, 19)
```

```
In [14]: df['ANNUAL'].value_counts()
```

```
Out[14]: ANNUAL
0.0      26
1024.6    4
790.5     4
770.3     4
1114.2    3
..
419.8     1
428.9     1
527.8     1
322.9     1
1642.9    1
Name: count, Length: 3713, dtype: int64
```

```
In [15]: df['Jan-Feb'].value_counts()
```

```
Out[15]: Jan-Feb
0.0      244
0.1       80
0.2       52
0.3       38
0.4       32
...
58.2       1
23.3       1
95.2       1
76.9       1
69.3       1
Name: count, Length: 1220, dtype: int64
```

```
In [16]: df['Mar-May'].value_counts()
```

```
Out[16]: Mar-May
0.0      38
0.1      13
0.3      11
8.3      11
2.7      10
..
249.5     1
148.8     1
191.9     1
207.0     1
223.9     1
Name: count, Length: 2262, dtype: int64
```

```
In [17]: df['Jun-Sep'].value_counts()
```


```
Out[17]: Jun-Sep
0.0      10
573.8     4
613.3     4
434.3     4
334.8     4
..
1328.5     1
1073.1     1
1373.0     1
897.7     1
958.5     1
Name: count, Length: 3684, dtype: int64
```

```
In [18]: df['Oct-Dec'].value_counts()
```

```
Out[18]: Oct-Dec
0.0      29
0.1      15
0.5      13
0.6      12
0.7      11
..
41.5     1
95.4     1
11.7     1
72.9     1
555.4     1
Name: count, Length: 2389, dtype: int64
```

## Exploratory Data Analysis

```
In [19]: df=df[['JAN', 'FEB', 'MAR', 'APR', 'DEC']]
sns.heatmap(df.corr(),annot=True)
plt.show()
```



```
In [20]: df.columns
```

```
Out[20]: Index(['JAN', 'FEB', 'MAR', 'APR', 'DEC'], dtype='object')
```

```
In [21]: x=df[["FEB"]]
y=df["JAN"]
```

## 1.Linear Regression

```
In [22]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [23]: from sklearn.linear_model import LinearRegression
reg=LinearRegression()
reg.fit(X_train,y_train)
print(reg.intercept_)
coeff_=pd.DataFrame(reg.coef_,x.columns,columns=['coefficient'])
coeff_
```

```
9.640361097385627
```

```
Out[23]:
```

	coefficient
FEB	0.442528


```
In [24]: score=reg.score(X_test,y_test)
print(score)
```

```
0.17932442210801225
```


```
In [25]: predictions=reg.predict(X_test)
```

```
In [26]: plt.scatter(y_test,predictions)
```

```
Out[26]: <matplotlib.collections.PathCollection at 0x172c148e590>
```




```
In [27]: df500=df[:][:500]  
sns.lmplot(x="FEB",y="JAN",order=2,ci=None,data=df500)  
plt.show()
```



```
In [28]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33)  
reg.fit(X_train,y_train)  
reg.fit(X_test,y_test)
```

```
Out[28]: LinearRegression  
LinearRegression()
```

```
In [29]: y_pred=reg.predict(X_test)
plt.scatter(X_test,y_test,color='black')
plt.plot(X_test,y_pred,color='red')
plt.show()
```



```
In [30]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(X_train,y_train)
y_pred=model.predict(X_test)
r2=r2_score(y_test,y_pred)
print("R2 Score:",r2)
```

R2 Score: 0.2575267236574985

```
In [31]: #Conclusion:This model has 10% of accuracy
```

## 2.Ridge Regression

```
In [32]: from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

```
In [33]: features= df.columns[0:5]
target= df.columns[-5]
```

```
In [34]: x=np.array(df['JAN']).reshape(-1,1)
y=np.array(df['FEB']).reshape(-1,2)
```

```
In [35]: x= df[features].values
y= df[target].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=17)
```

```
In [36]: ridgeReg=Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
train_score_ridge=ridgeReg.score(x_train,y_train)
test_score_ridge=ridgeReg.score(x_test,y_test)
```



```
In [37]: print("\n Ridge Model:\n")
print("the train score for ridge model is{}".format(train_score_ridge))
print("the test score for ridge model is{}".format(test_score_ridge))
```


Ridge Model:

the train score for ridge model is 0.9999999999874308  
the test score for ridge model is 0.9999999999883638

```
#conclusion:
the train score for ridge model is 0.9999999999874192
the test score for ridge model is 0.99999999998833
```

```
In [38]: lr=LinearRegression()
```

```
In [41]: ze= (10,10))
s,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=5,color="red",label=r'Ridge;$alpha=10$',zorder=7)
s,alpha=0.4,linestyle='none',marker='o',markersize=7,color="green",label='Linear Regerssion')
ion = 90)
```



### 3.Lasso Regression


```
In [43]: print("\n Lasso Model:\n")
lasso=Lasso(alpha=10)
lasso.fit(x_train,y_train)
train_score_ls=lasso.score(x_train,y_train)
test_score_ls=lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.9999207503194595  
The test score for ls model is 0.9999206588980594

```
In [44]: pd.Series(lasso.coef_,features).sort_values(ascending=True).plot(kind="bar")
```


Out[44]: <Axes: >



```
In [45]: from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,1,10],random_state=0).fit(x_train,y_train)
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

0.9999999999999921  
0.9999999999999921

```
In [47]: e= (10,10))
,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=5,color='red',label=r'Ridge;\alpha=10$',zorder=7)
.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'Ridge;\alpha-grid$')
,alpha=0.4,linestyle='none',marker='o',markersize=7,color="green",label='Linear Regression')
on = 90)
```



## 4.Elastic Net

```
In [48]: from sklearn.linear_model import ElasticNet
elnet=ElasticNet()
elnet.fit(x,y)
print(elnet.coef_)
print(elnet.intercept_)
print(elnet.score(x,y))

[9.99098882e-01 0.00000000e+00 2.95025317e-05 0.00000000e+00
 0.00000000e+00]
0.016260183535354855
0.9999992158680019
```

```
In [50]: y_pred_elastic = eln.predict(x_train)
mean_squared_error=np.mean((y_pred_elastic - y_train)**2)
print(mean_squared_error)

0.0008817707583048099
```

## Conclusion

THE SCORE OF LINEAR REGRESSION IS :- 0.1793580786264921  
THE SCORE OF RIDGE MODEL IS :- 0.99999999998833  
THE SCORE OF LASSO MODEL IS :- 0.99999999999992  
THE SCORE OF ELASTIC NET IS :- 0.9999992160905338  
\*AMONG ALL MODELS LASSO YEILD HIGHEST ACCURACY.SO,WE PREFER LASSO MODEL FOR THIS DATA SET\*

```
In [ ]:
```